

Blackjack++

making counterfactual regret minimization more accessible

Mateo Maturana
Marcus Fong
Erin Liang

A green felt casino table, likely for blackjack, is shown with various playing cards and stacks of multi-colored chips. The cards include a King of Diamonds, Queen of Hearts, Jack of Clubs, and several other cards. The chips are stacked in several locations. The text 'PERFECT PAIRS' is visible on the table layout. The background is slightly blurred, focusing attention on the game elements.

Intro

Motivation ++ Goals ++ Background Information

Motivation

- Algorithmic basis for the bots dominating annual poker competitions is **counterfactual regret minimization (CFR)** [1]
- Regret-based algorithms are nascent (~2014 for CFR+)
 - **Few materials available** to introduce students, researchers, practitioners
 - Best way to learn game strategy + CFR is to see how AI's optimal strategy iterates throughout a game

Solving Large Imperfect Information Games
Using CFR⁺

Oskari Tammelin
ot@iki.fi

July 21, 2014

Abstract

Counterfactual Regret Minimization and variants (e.g. Public Chance Sampling CFR and Pure CFR) have been known as the best approaches

Motivation

- Why C++?
 - CFR Games → large game trees
 - CFR requires a lot of iterations to converge
 - Object-Oriented → conceptually intuitive to understand
 - modern C++ opportunities: coroutines and modules
- Applying game theory class learnings! 🎉

Goals

- Design CFR decision tree class using the algorithm described in Tammelin's "Solving Large Imperfect Information Games Using CFR+"
- Need an easy-to-understand game to apply CFR to → Blackjack
 - Blackjack is a game of imperfect knowledge, which makes it the perfect game to apply CFR
- We want modify an existing game to ensure that we don't repeat work that has already been done
- Design the game using modern C++ techniques

CFR Overview

- Recursive algorithm that attempts to minimize regret at each iteration vs our own current strategy
- Strategy at each iteration is computed using RegretMatching
 - Definition: Regret for Action A is defined as the amount we “regret” not doing this action. Mathematically, this is $\text{Utility}[A] - \text{Utility}[\text{current mixed strategy}]$
 - Strategies are then computed directly proportional to regret
 - Ex: The current utility of our strategy is 3. Utility of A is 7, Utility of B is 5, Utility of C is 1. Therefore the regret of A is 4, B is 2, and C is 0 (regret has a floor of 0). Thus, our strategy is the next iteration will be to do A with $\frac{2}{3}$ frequency, B with $\frac{1}{3}$ frequency, and C with 0 frequency.
- Final strategies are taken by averaging the strategy at each iteration
- Overtime, the average strategy will converge to nash equilibrium

CFR Overview: CFR+

- In order to reach faster convergence, we use CFR+ in our implementation.
- RegretMatching+
 - Instead of computing the regret at each individual iteration, we keep a sum of all regrets so far
 - Strategy is computed proportionally to the regret **sums**
- Weighted Strategy Sums
 - Recall that in CFR, the final strategy is computed using the average strategy from each iteration
 - However, shouldn't strategies in later iterations be closer to equilibrium?
 - Weighted Strategy Sums weight later iterations more for the final strategy
 - $\text{strategySum}[i] += \text{curStrategy}[i] * \text{iterationNum}$
 - Sums are normalized after all iterations are complete
- CFR+ can help reach convergence orders of magnitude faster

Blackjack++ Overview and Rules

- All players and the dealer start with 1 card face up, 1 card face down.
- Players take turns choosing to either **hit/stand**. Once both players stand, their action is over
 - **Hit:** A random card is drawn and given to the player
 - **Stand:** The player indicates they no longer want any cards
- After player action is over, the dealer draws their cards according to an established protocol (hit if <17 , stand if ≥ 17).
- The player/dealer with the highest sum is awarded the pot, Ace can be either 1 or 11.

The background is a blurred, high-angle shot of a poker table. It shows several playing cards scattered across the green felt, including a King of Diamonds, a Queen of Hearts, and a Jack of Clubs. There are also stacks of colorful poker chips in the upper left and lower right corners. The text 'PERFECT PAIRS' is visible on the table layout. The overall image has a soft, out-of-focus aesthetic.

Tutorial

How to Build and Run ++ “Hello World” ++ Common Errors

Usage: How to Build and Run

We included a Makefile to make building Blackjack++ easier:

```
$ make
```

```
c++ -I/usr/local/opt/ncurses/include -c -o src/Cfr.o src/Cfr.cpp  
[...]
```

To run a game of Blackjack++ to save under `sol` where all players start with a 10 with a maximum of 1 hit (players can draw up to 1 more card) and 10000 iterations of CFR:

```
$ ./a.out "10, 10, 10" 1 10000 sol
```

In the general case, we would run as follows:

```
$ ./a.out "[player0 card], [player 1 card], [dealer card]" [maxHits]  
[numIterations] [dump directory]
```

Blackjack++ Game Example

- How does our version of Blackjack work...
- Game start:
 - My deck: {5 (d), 3 (u)}
 - Other deck: {2 (d), 10 (u)]
 - I can only see the upcards of other players
 - I Hit and draw a 8: {5 (d), 3 (u), 8 (u)}
 - Because I saw that he has a 10 and I have an 8, I might as well hit because I never bust at this point level
 - P2 hits and draws a 8
 - Given that the other player hasn't busted, I know that the down card has to be a small number. I know that P2 has at least 2 more than me, so I might as well take the chance to win.
 - I Hit and draw a 5, putting me at 21 and a win

Blackjack++ Game Example Output (Directory Structure)

```
game1/ -  
  | - strategy.txt  
  | - s  
    | - strategy.txt  
    | - s  
    ...  
  | - h1  
    | ...  
  | - h2  
    | ...  
  | - h3  
    | ...  
  ...
```

master		cpp_blackjack / game1 /	Go to file	Add file	...
maturanamateo dumping files 99d0858 1 hour ago History					
..					
h1	dumping files	1 hour ago			
h10	dumping files	1 hour ago			
h2	dumping files	1 hour ago			
h3	dumping files	1 hour ago			
h4	dumping files	1 hour ago			
h5	dumping files	1 hour ago			
h6	dumping files	1 hour ago			
h7	dumping files	1 hour ago			
h8	dumping files	1 hour ago			
h9	dumping files	1 hour ago			
s	dumping files	1 hour ago			
strategy.txt	dumping files	1 hour ago			

Blackjack++ Game Example Output (Strategy File)

Current Acting: 1

```
Down Card: 1 || Stand: 0.000516931, Hit: 0.999483
Down Card: 2 || Stand: 0.000396455, Hit: 0.999604
Down Card: 3 || Stand: 5.89984e-05, Hit: 0.999941
Down Card: 4 || Stand: 1.70004e-05, Hit: 0.999983
Down Card: 5 || Stand: 4.04662e-05, Hit: 0.99996
Down Card: 6 || Stand: 0.0182729, Hit: 0.981727
Down Card: 7 || Stand: 0.055304, Hit: 0.944696
Down Card: 8 || Stand: 0.0703224, Hit: 0.929678
Down Card: 9 || Stand: 0.75498, Hit: 0.24502
Down Card: 10 || Stand: 0.990737, Hit: 0.00926331
```

Common Errors

Mostly errors with command-line arguments:

1. Setting `maxHits` too high

- maximum number of times that a player can request an extra card
- would cause the game tree build time to be **very high** ... foreshadowing for performance
- Why? branching factor of ~100 for an increment of 1 in `maxHit`

2. Setting `cfrIterations` too high

- number of iterations user wants CFR to take
- due to CFR+, strategies converge fairly quickly

The background is a blurred, high-angle shot of a poker table. It shows several playing cards scattered across the green felt, including a 4 of hearts, 6 of diamonds, 7 of clubs, and 9 of spades. There are also stacks of multi-colored chips (red, yellow, green, blue) and betting areas labeled "PERFECT PAIRS".

Design Manual

CFR ++ Game ++ Player ++ State Overview
(more details in the design doc)

CFR Class

- Runs all CFR relevant code
- `train()` – Runs all CFR iterations. For each iteration, generates a set of down cards
- `normalize()` – Normalizes strategy sums after all CFR iterations have finished
- `runCfr()` – Computes utilities and regrets
- `getStrategy()` – Uses RegretMatching+ to get the current strategy
- `getTerminalNodePayoffs()` – Generates payoff matrix for any terminal node
- All objects taken by reference. They are produced in tree creation, so want to avoid making unnecessary copies.

Game Class

- Defines the logic for the Blackjack++ game
- Starting from `main()` ...
 - Game object created using command line args: `startingCards`, `maxHits`, `cfrIterations`
 - `constructTree()` — Builds game tree
 - Starts off larger Blackjack++ game logic
 - Keeps track of time and amount of nodes it takes to build the game tree for performance
- Other methods for debugging + pedagogical purposes
 - `printRandomPath()`
 - `dumpToFiles()`
- Alternative approaches?

Player Class

- Two Players, one dealer.
- What do we need to keep track of about the Player during Blackjack++?
 - Which player is this?
 - `id`
 - Is the player allowed to take any more actions?
 - `doneAction`
 - When is a player done?
 - `cards sum > 20` (no feasible way they can take more cards because their minimum sum is 21)
- Constructed in Game class in `Game::constructTree()`

```
auto Player0 = Player(0, false);
```

```
auto Player1 = Player(1, false);
```

State Class

- What is a State? What does it represent?
- Where is our class used with respect to CFR algorithm?
- `State::State()`
- `State::populateChildren()`
- `State::createStandState()`
- `State::createHitState()`

```
std::vector<std::future<State>> futures;
futures.reserve(11);

futures[0] = std::async(std::launch::async, [this]() {
    return State::createStandState();
});
for (int i = 1; i <= 10; i++) {
    futures[i] = std::async(std::launch::async, [this, i]() {
        return State::createHitState(i);
    });
}

for (int i = 0; i <= 10; i++) {
    if (i == 0) {
        const auto& state = futures[0].get();
        //state.printDetails();
        nextStates.emplace("s", state);
    } else {
        std::string formatted = "h" + std::to_string(i);
        const auto& state = futures[i].get();
        //state.printDetails();
        nextStates.emplace(formatted, state);
    }
}
```


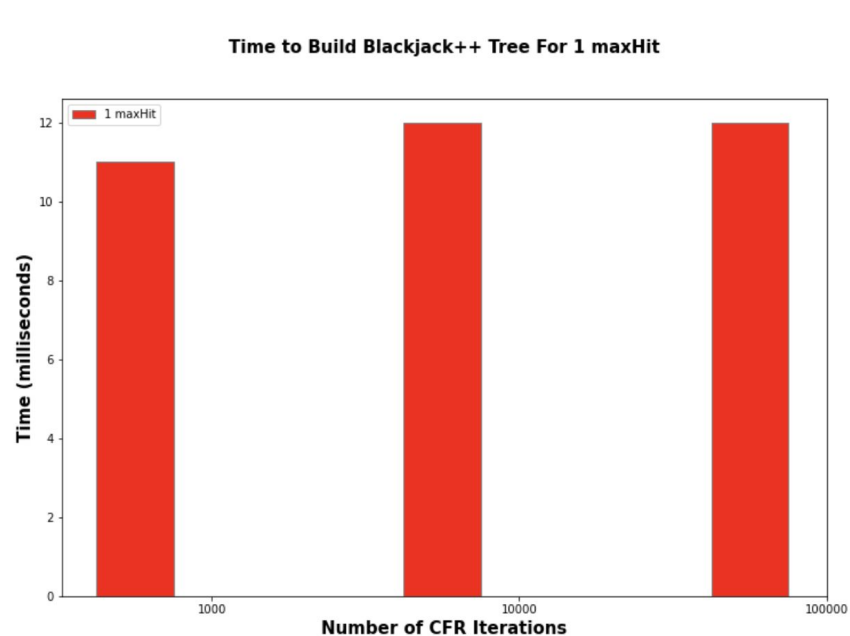
The background is a blurred, high-angle shot of a green poker table. Several playing cards are scattered across the table, including a King of Diamonds, a Queen of Hearts, a Jack of Clubs, and a 10 of Spades. There are also stacks of colorful poker chips in the upper left and lower right corners. The text 'PERFECT PAIRS' is visible on the table layout. The overall image has a soft, out-of-focus aesthetic.

Performance

some pretty graphs!



Performance

Impact of $maxHits$ and $cfrIterations$ on Game Tree Build Time



Performance


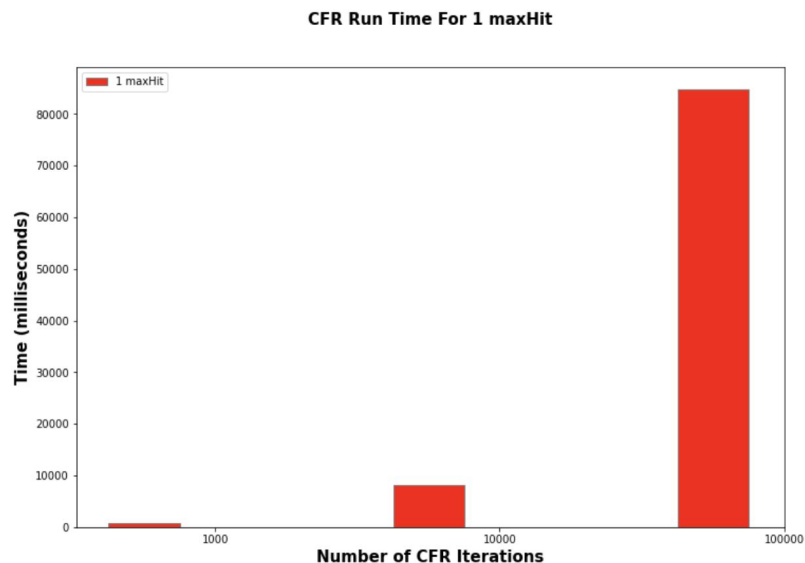
Impact of maxHits and cfrIterations on Game Tree Build Time



- factor of ~ 100 increase in build time when maxHits incremented \rightarrow branching factor
- no impact on build time when cfrIterations increased by factor of 10


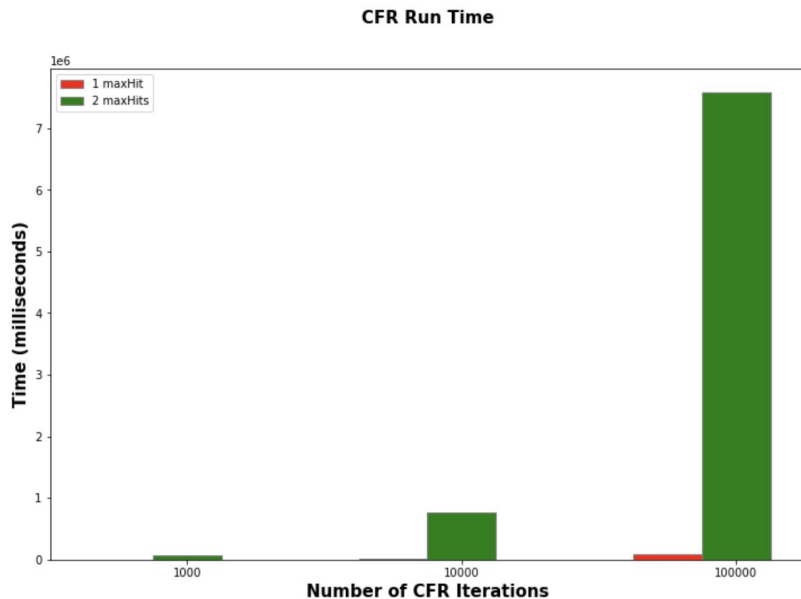
Performance

Impact of maxHits and cfrIterations on CFR Training Time



Performance

Impact of maxHits and cfrIterations on CFR Training Time



- factor of ~ 100 increase in CFR training time when maxHits incremented \rightarrow branching factor
- factor of ~ 10 increase in CFR training time when cfrIterations increased by factor of 10

The background is a blurred, high-angle shot of a poker table. It shows various playing cards scattered across the green felt, including a 10 of hearts, a 9 of hearts, a 10 of diamonds, a 9 of diamonds, a 10 of clubs, a 9 of clubs, a 10 of spades, and a 9 of spades. There are also several stacks of colorful poker chips (red, yellow, green, blue) and a small pile of chips. The text "PERFECT PAIRS" is visible on the table layout. The overall image has a soft, out-of-focus aesthetic.

Post-Mortem

Reflections ++ Roadblocks ++ Future Work

Reflections + Roadblocks

- Current reflections
 - Positives
 - Negatives
- Integrating with modern C++
 - Coroutines
 - Modules

Future Work

- Interactivity
 - Being able to play vs the computed strategies
 - In theory, over the long run the computed strategies should not lose
- Generalizing to other Games
 - Building a CFR library that can build game trees for other types of games easily
- UI
 - Getting the current strategies requires moving through a lot of directories
 - In an ideal world, we would want players to easily find strategies for decision points
 - For example, a lot of online poker solvers have really good UIs (see next slide)

CASH

6max

NL50

General

100 bb

UTG

100

FOLD

RAISE 2

ALLIN 100

HJ

100

FOLD

RAISE 2

ALLIN 100

CO

100

FOLD

RAISE 2.3

ALLIN 100

BTN

100

FOLD

RAISE 2.5

ALLIN 100

SB

99.5

FOLD

CALL

RAISE 12

ALLIN 100

BB

99

FOLD

CALL

RAISE 13

ALLIN 100

FLOP

5.5

J

T

7

BB

97.5

CHECK

BET 1.8 (33%)

BTN

97.5

CHECK

BET 1.8 (33%)

BET 2.75 (50%)

BET 4.1 (75%)

BET 6.9 (125%)

STRATEGY

RANGES

BREAKDOWN

REPORTS

AA	AKs	AQs	AJs	ATs	A9s	A8s	A7s	A6s	A5s	A4s	A3s	A2s
AKo	KK	KQs	KJs	KTs	K9s	K8s	K7s	K6s	K5s	K4s	K3s	K2s
AQo	KQo	QQ	QJs	QTs	Q9s	Q8s	Q7s	Q6s	Q5s	Q4s	Q3s	Q2s
AJo	KJo	QJo	JJ	JTs	J9s	J8s	J7s	J6s	J5s	J4s	J3s	J2s
ATo	KTo	QTo	JTo	TT	T9s	T8s	T7s	T6s	T5s	T4s	T3s	T2s
A9o	K9o	Q9o	J9o	T9o	99	98s	97s	96s	95s	94s	93s	92s
A8o	K8o	Q8o	J8o	T8o	98o	88	87s	86s	85s	84s	83s	82s
A7o	K7o	Q7o	J7o	T7o	97o	87o	77	76s	75s	74s	73s	72s
A6o	K6o	Q6o	J6o	T6o	96o	86o	76o	66	65s	64s	63s	62s
A5o	K5o	Q5o	J5o	T5o	95o	85o	75o	65o	55	54s	53s	52s
A4o	K4o	Q4o	J4o	T4o	94o	84o	74o	64o	54o	44	43s	42s
A3o	K3o	Q3o	J3o	T3o	93o	83o	73o	63o	53o	43o	33	32s
A2o	K2o	Q2o	J2o	T2o	92o	82o	72o	62o	52o	42o	32o	22

UTG

HJ

CO

BTN

SB

BB

5.5 BB

J

T

7

Bet 6.9

125%

6.2%

30.24

combos

Bet 4.1

75%

15.8%

76.79

combos

Bet 2.75

50%

10.8%

52.64

combos

Bet 1.8

33%

9.7%

47.23

combos

Check

57.4%

278.82

combos

HANDS

SUMMARY

FILTERS

BLOCKERS

HANDS

Straight

1%

Set

1.9%

Two pair

2.7%

Overpair

3.7%

Top pair

12.5%

Second pair

12.4%

Third pair

9%

Low pair

5.8%

Ace high

26.6%

King high

13.3%

No made hand

11.3%

EQUITY BUCKETS

References

1. [Solving Large Imperfect Information Games Using CFR+ Paper](#)

The background is a blurred, high-angle shot of a green casino table. Several playing cards are scattered across the surface, including a 10 of hearts, a 9 of diamonds, a 7 of clubs, and a 6 of spades. Stacks of colorful chips (red, yellow, green) are visible in the upper left and lower right corners. The text "PERFECT PAIRS" is printed on the table layout in several locations. The overall image has a soft, out-of-focus aesthetic.

Thanks for a great semester!

Any questions?